

MA22019 Introduction to Data Science - Lecture Notes

Dr Christian Rohrbeck

2025-02-01

Table of contents

Chapter 1

MA22019 Introduction to Data Science - Lecture Notes

Overview

Content

In practice we often have to deal with large and complex data sets (“**big data**”) - many organisations and companies have collected vast amounts of data over the last decade. Analyzing such data and extracting valuable insights from it lies at the heart of data science.

Although there exists no unique definition, the consensus is that data science is interdisciplinary and requires combining expertise across mathematics, computer science and the area of application. As such, we as mathematicians need to have at least some understanding of the context in which the data were collected / are analyzed in order to extract the important information and to communicate effectively with other disciplines.

In this course we will focus on four key areas of data science:

- Data wrangling
- Data visualization
- Text data analysis
- Spatial data analysis

Throughout this unit we will usually work with real-world data sets which you can download from the MA22019 Moodle page.

Aims and Objectives

After taking this unit, you should be able to:

- Demonstrate knowledge of data science techniques used for data wrangling and visualization.
- Use R for data visualization and wrangling.
- Show awareness of the applications of these methods.
- Understand some of the mathematical models underlying spatial and text data analysis.
- Apply these methods in R to analyse large and complex data sets. This includes demonstrating a practical and critical approach to data analysis, including the ability to select a suitable data science method, explaining your reason for choosing it, and correctly interpreting the results.

Prerequisites

You should be familiar with the basics of R introduced in the Probability and Statistics component of MA12003. On Moodle, you can find the “Brief introduction to R” provided by Professor Jennison in case you are unsure about certain aspects.

This course will also use some of the fundamentals introduced in Year 1, including the various probability distributions, and the concepts of mean, variance and correlation.

Assessment

Summative Assessment

Your mark for MA22019 is based on two individual pieces of coursework:

- **Coursework 1: 40% of unit mark.** Set at the end of Week 3 and due at the end of Week 4. This coursework will focus on data wrangling and data visualization.
- **Coursework 2: 60% of unit mark.** Set at the end of Week 10 and due at the start of Revision Week. This coursework will focus on text and spatial data analysis, but you will also have to demonstrate your ability to use data wrangling and data visualization techniques.

You have to submit your solutions to the coursework via Moodle.

Formative Assessment

Problem sheets will be set at the beginning of the week (except for Weeks 4, 10 and 11) and include:

Exercises: Focus more on the programming aspects of the course and help you to revise the content covered in that week’s lectures. You can submit your solutions to most questions via a Moodle quiz and you will receive direct feedback.

Tutorial questions: To be attempted in the weekly tutorials. Solutions will be made available on Tuesday evening after the last tutorial has finished.

Homework question: Open-ended question, similar in style to the coursework questions. Your answer can be submitted via Moodle to your tutor for feedback. After the submission deadline, I will provide a solution which demonstrates the aspects one may consider to address the question (and achieve a first-class mark), but I won’t provide a fully formulated answer - I want you to practice discussing your results in your own words.

Organisation

Lecture Notes

Lecture notes are available from the MA22019 Moodle page. The notes are available in two formats (HTML and PDF) with identical content but different layouts.

Important: While the lecture notes and problem class notes provide most of the relevant information on the mathematics/statistics and programming aspects of the course, they do not fully explore (1) how to decide which plots are best suited for addressing a research question and (2) how to interpret these plots in the context of the application. These skills are best trained by looking at a range of data applications, and we will do so in the lectures and problem classes. The open-ended question on the problem sheets is designed to support your learning and to provide additional practice on this topic. As such, you are strongly encouraged to actively engage with the lectures and problem sheets.

Lectures and Tutorials

In person lecture: Wednesday 10:15-12:05 (Weeks 1-3,5-9) or 10:15-11:05 (Weeks 4,10,11) in 3WN 2.1

Live online learning session (LOIL): Friday 14:15-15:05 (Weeks 1-3) on Zoom (link on Moodle)

Computer Lab: You will be assigned to a small group which will meet weekly to go over the tutorial questions on the problem sheet. Please check your timetable to identify the time and location of your tutorial.

Schedule

The rough outline for the individual weeks is as follows:

	Week	Topic	Remarks
1		Data Wrangling	Tutorial on RMarkdown
2		Data Wrangling + Data Visualization	Problem Sheet 1
3		Data Visualization	Problem Sheet 2
			Release of Coursework 1
4		Q&A Coursework 1	Deadline Coursework 1
5		Text Data Analysis 1	No tutorials
6		Text Data Analysis 2	Problem Sheet 3
7		Spatial Data Analysis 1	Problem Sheet 4
8		Spatial Data Analysis 2	Problem Sheet 5
9		Spatial Data Analysis 3	Problem Sheet 6
10		Revision Class	Release of Coursework 2
11		Q&A Coursework 2	
12			Deadline Coursework 2

We will finish the core content by the end of Week 9. In Weeks 4, 10 and 11, we will only have Revision and Coursework Q&A sessions.

R and RStudio

We will use R and RStudio throughout this course. You can access RStudio via the UniDesk, but I would recommend that you install R and RStudio on your own computer / laptop. If you already have R installed, please check that you have at least version 4.4.0.

To install all the relevant packages, run the R code in the file “InstallPackages.R”, which you can find on Moodle. The individual packages are then loaded using the `library()` function, e.g.

```
library(dplyr)
```

This has to be done every time you start R/RStudio.

To load data from external files, we will use the following two functions

```
load( "..." )      # To load .RData files
read.csv( "..." ) # To load .csv files
```

Resources

This unit is self-contained in the sense that you will not need to read text books. However, you may wish to consult the following books, available as ebooks via the University Library, to support your learning and understanding:

Boehmke, Bradley C. *Data Wrangling with R*. 1st Edition. 2016. Springer.

Mailund, Thomas. *Beginning Data Science in R : Data Analysis, Visualization, and Modelling for the Data Scientist*. 1st Edition 2017. Apress.

Wickham, Hadley. *ggplot2 : Elegant Graphics for Data Analysis*. 2nd Edition. 2016. Springer.

Silge, Julia and Robinson, David. *Text Mining with R: A Tidy Approach*. 2015. O'Reilly.

Remark: Some examples in these books use R functions that have been superseded, i.e., it is advised to use an alternative, more recent function (which may have a different syntax). I will use the most recent functions in this course.

Chapter 2

Data Wrangling

In practice, **raw data** (measurements, etc.) may be provided in a format that hampers its analysis. Common problems include:

- Data types are incompatible with available R functions - data may, for instance, be imported as a set of characters, but we require numerical values for most R functions;
- Variable names are uninformative - we do not want to always go back and look up what the individual variables represent;
- We are only interested in a subset of the data - a study may have produced a lot of data, but we want to focus on a specific aspect.

Data wrangling is concerned with restructuring the raw data into a format more useful for the analysis. We also want to extract key information from the data, i.e, performing some **data exploration**.

In this chapter we will explore how we can perform such tasks in R. Most of the functions we will use are provided by the **dplyr** R package (and other R packages where appropriate). Let's load the dplyr package,

```
library( dplyr )
```

When we import data into R, the data will usually be stored as a **data frame**, which corresponds to a matrix, where each column has a name. The dplyr package is particularly effective at working with this data format. To access the individual columns for the different variables in the data frame, we can use the \$ sign, followed by the name of the column (you will see this syntax at multiple points).

2.1 Data cleaning

After loading the data, the first step is to check that variables have the correct data type and decide whether the variable names are suitable/informative. If the data type is incorrect, we should convert it, in particular if we want to apply R functions to the data. The most common conversion we will have to make is from the type **character** to the type **numeric** or **date**. In this section, we highlight some of the available R functions for converting and renaming variables, and illustrate their application using publicly available river flow data.

2.1.1 Converting characters into numerical values

When loading the data for a numerical variable into R, the individual values may be stored as strings/words. Possible reasons include, for instance, that the value 2345.34 is stored as “2,345.34” in the data file (see Problem Class 1) or that missing values are represented via a letter - once R fails to convert a single entry to a numerical value, the whole column (variable) is converted to type **character**. However, this data type is often not useful because only a few functions can work with it.

We will introduce the functions **as.numeric()** and **case_when()** that may be used to convert a character to a numerical value.

Example 1: Suppose we work with a data set where the letter “M” is used to indicate that an observation is missing. A toy example is provided in the data file “DataCleaningExample1.csv”. We load this data and use the function **glimpse()** in the dplyr package to print the data type and values:

```
Example1 <- read.csv( "Data/DataCleaningExample1.csv" )
glimpse( Example1 )
```

```
Rows: 8
Columns: 1
$ Value <chr> "1.02", "0.98", "0.79", "M", "2.1", "15.1", "M", "4.2"
```

We find that the values are stored as characters/words, as indicated by the data type being **<chr>**. In such a situation, we cannot use the **mean()** function to derive the average value - there is no average of a set of words:

```
mean( Example1$Value )
```

```
Warning in mean.default(Example1$Value): argument is not numeric or logical:
returning NA
```

```
[1] NA
```

Instead, we have to first use the function **as.numeric()** to convert the words to numerical values

```
Example1$Value <- as.numeric( Example1$Value )
```

```
Warning: NAs introduced by coercion
```

```
glimpse( Example1 )
```

```
Rows: 8
Columns: 1
$ Value <dbl> 1.02, 0.98, 0.79, NA, 2.1, 15.1, NA, 4.20
```

We see that the data type has changed to **<dbl>**, which is one data type for numerical values. Further, the R output shows that all entries with the letter “M” were converted to NA (not available) - this is R’s way to tell us that the conversion did not work or that a value is missing (which is exactly what we want here).

With the values converted to the correct type, we can now calculate their mean:

```
mean( Example1$Value, na.rm=TRUE ) # na.rm=TRUE to ignore the entries with NA
```

[1] 4.031667

Example 2: Suppose responses to a survey question were encoded as “Y” (“Yes”) or “N” (“No”) and we received the following data vector:

```
responses <- c("Y", "Y", "N", "Y", "N", "Y", "N", "N", "N", "Y")
```

Again, we cannot use `mean(responses)` to derive the proportion of participants who answered with “Yes”, because the `mean()` function requires numerical or logical values.

One common approach to derive the proportion in practice is to encode the outcomes as numerical values to which the `mean()` function is then applied. The `dplyr` package provides the function `case_when()` which allows us to replace “Y” and “N” by 1 and 0 respectively:

```
responses <- case_when( responses == "Y" ~ 1, responses == "N" ~ 0 )
mean( responses )
```

[1] 0.5

We find that 50% of the participants answered the question with “Yes”. One strength of `case_when()` is that we can define as many cases as we need, there is no limit. The function can also be used to convert other types of data, and it is not limited to converting a character into a numerical value.

Remark 1: If you forget to specify a case in `case_when()`, the converted value for any unspecified case will be `NA` by default. The default option can be changed and for our example we could have also used

```
responses <- case_when( responses == "Y" ~ 1, .default = 0 )
```

Remark 2: Be aware that `case_when()` considers the expressions sequentially (just as when you are using if, else if and else statements). The following pieces of code show an example where the result depends on the order of the conditions:

```
x <- c( 10, 20, 40 )
case_when( x %% 10 == 0 ~ 1, x %% 20 == 0 ~ 2 )
```

[1] 1 1 1

We see that all converted values are all equal to 1. Let’s see what happens when we change the order of the conditions:

```
case_when( x %% 20 == 0 ~ 2, x %% 10 == 0 ~ 1 )
```

[1] 1 2 2

This second result seems more intuitive. Consequently, we should proceed from the most specific to the most general condition when using `case_when()`.

2.1.2 Converting characters into dates

In many studies we are provided with the time the data were observed. This information is often important in applications and we cannot simply ignore it. When loading variables representing dates into R, their values are often stored as strings, such as “01/10/2022”.

The R package **lubridate** provides a range of nice functions to convert data of type **character** into the data type **date** or **date-time**. For instance, to convert the character expressions “01/10/2022” and “15/10/2023”, we use the **as_date()** function,

```
library(lubridate)
date_observed <- c( "01/10/2022", "15/10/2023" )
date_converted <- as_date( date_observed, format="%d/%m/%Y" )
glimpse( date_converted )
```

```
Date[1:2], format: "2022-10-01" "2023-10-15"
```

We see that the default output format for dates is **year-month-day**.

Remark 1: After converting values to **date**, we can extract the year and month using the functions **year()** and **month()** respectively. Let’s extract the year from the dates in **date_converted**:

```
year( date_converted )
```

```
[1] 2022 2023
```

Remark 2: We can also calculate the difference between dates. For instance, if we consider the two entries in the vector of converted dates, we find

```
date_converted[2] - date_converted[1]
```

```
Time difference of 379 days
```

So you can now use R to quickly calculate how many days there are left until the Easter break.

2.1.3 Changing variable names

We should avoid using uninformative (or very long) variable names. Let’s generate a data frame with two columns, where each column contains five samples from a standard normal distribution

```
set.seed( 2025 )
obs <- data.frame( "x"=rnorm(5, mean=0, sd=1), "y"=rnorm(5, mean=0, sd=1) )
obs
```

	x	y
1	0.6207567	-0.16285434
2	0.0356414	0.39711189
3	0.7731545	-0.07998932
4	1.2724891	-0.34496518
5	0.3709754	0.70215136

We may argue that the variable names **x** and **y** are uninformative and should be changed to **Sample1** and **Sample2** respectively. The function **rename()** in the **dplyr** R package allows us to do this:

```
obs <- rename( obs, "Sample1"=x, "Sample2"=y )
obs
```

	Sample1	Sample2
1	0.6207567	-0.16285434
2	0.0356414	0.39711189
3	0.7731545	-0.07998932
4	1.2724891	-0.34496518
5	0.3709754	0.70215136

2.1.4 Example: Loading and cleaning NRFA river flow data

The National River Flow Archive (www.nrfa.ceh.ac.uk) provides data for hundreds of sites (gauges) across the UK. We want to analyze daily river flow data for the River Avon at Bathford. The data are available in the file “Bathford River Flow.csv”.

When looking at the data file, we identify two aspects that need to be taken into account when loading the data into R

- 1) The first 20 lines are data descriptors (so called **meta data**), while the remaining lines contain the actual data: dates and river flow measurements.
- 2) The letter “M” appears in the third column whenever the river flow measurement is missing in later years.

To ignore the first 20 lines and avoid importing the data file in a wrong format, we have to use three of the options provided by the `read.csv()` function:

```
Bathford_RF <- read.csv( "Data/Bathford River Flow.csv", skip=20, header=FALSE,
                           colClasses = c("character","numeric","NULL") )
```

```
Warning in read.table(file = file, header = header, sep = sep, quote = quote, :
  cols = 2 != length(data) = 3
```

The option `skip=20` means we ignore the first 20 lines, while `colClasses= c("character","numeric","NULL")` leads to the third column being ignored when loading the data (you can ignore the warning message in this case). Finally, we set `header=FALSE`, because the file does not provide variable names.

Tip: Have a look at the data set in the data file before trying to load it. R (in particular recent versions) may load the data into an incorrect format instead of giving an error. As an example, remove the option `colClasses=...`. You will find that the number of observations increases, but some of the dates are now listed as “M”.

Let’s have a look at the imported data using the `glimpse()` function,

```
glimpse( Bathford_RF )
```

```
Rows: 19,697
Columns: 2
$ V1 <chr> "1969-10-27", "1969-10-28", "1969-10-29", "1969-10-30", "1969-10-31~
$ V2 <dbl> 3.998, 3.958, 4.210, 4.480, 4.205, 3.830, 3.723, 3.986, 4.353, 4.52~
```

We see that there are 19,697 measurements in the data set. However, the variable names have to be changed - **V1** and **V2** are just not sensible. Further, we have to convert the dates into the **date** format. The river flow measurements are already stored as numeric values, so no conversion is required.

Let's start by changing the variable names to **Date** and **RiverFlow**

```
Bathford_RF <- rename( Bathford_RF, Date = V1, RiverFlow = V2 )
```

before converting the variable **Date** to the correct type,

```
Bathford_RF$Date <- as_date( Bathford_RF$Date, format="%Y-%m-%d" )
```

With the data having been **cleaned**, i.e., they have the correct type and informative names, we can start the analysis. As a first step, it is good practice to report the proportion of missing data of a variable. We can extract the proportion of missing river flow measurements using the functions **mean()** and **is.na()**

```
mean( is.na( Bathford_RF$RiverFlow ) )
```

```
[1] 0.006244606
```

We find that river flow measurements are missing on about 0.62% of dates and this should be reported.

Remark: Missing data is important when building models. In this course, you are only expected to state the proportion of missing data. The handling of missing data will be considered in more detail in the Year 3 unit MA32022 Statistical Modelling and Data Analytics 3A.

Now that the data frame is in a much better format, we can plot it using the function **plot()** covered in Year 1 Probability & Statistics:

```
plot( Bathford_RF$Date, Bathford_RF$RiverFlow, type='l',
      xlab="Date", ylab="River Flow", cex.lab = 1.5 )
```




Figure 2.1: River flow measurements at Bathford for 27 October 1969 to 30 September 2023.

Which of the following conclusions should we report when asked to comment on the frequency of river flow levels above $100\text{m}^3/\text{s}$ and the magnitude of river flow levels?

- 1) Recorded river flow levels were as high as approximately $250\text{m}^3/\text{s}$.
- 2) The data exhibits seasonality, with river flow levels being higher in winter than in summer.
- 3) The data covers the years 1969 to 2023.
- 4) There is at least one day with river flow levels exceeding $100\text{m}^3/\text{s}$ for most years.

2.2 Working with a single data frame

We now study a range of aspects that frequently come up during the data wrangling process:

Selecting subsets of observations and variables: When working with a large data set, not all variables and observations may be relevant. So we may want to reduce the size of our data set and only keep the observations required for our analysis. This process can happen before or after the data cleaning.

Deriving new variables from existing data: It may be useful to create new variables which we believe to be interesting to explore in our analysis. These new variables should be stored in the same data frame as the other variables.

Summarizing the data: As one of the first steps in the data exploration, we should derive summaries of the different variables to gain a better understanding of the data. For instance, as highlighted in Section 1.1.4, the proportion of missing data is one useful summary.

Sorting the data: In applications, interest may lie in extracting the smallest/largest observations and/or providing a ranking. As such, we need to be able to sort observations based on one, or more, criteria.

In this section we explore how to perform these operations using the dplyr package. In Problem Class 1 we will use the considered techniques to analyze a relatively large data set from Brazil. Other aspects of data exploration will be discussed in the next chapters.

Tip: All steps of the data wrangling / exploration process should be placed in an R (or R Markdown) script, so that we can make modifications quickly if something needs to be changed. It's also good practice to keep the raw data available in your R Workspace. In the following examples we never replace the raw data.

2.2.1 Filtering observations

In an analysis we may only want to focus on a subset of the data. For instance, when modelling the risk of flooding, we are mostly interested in the extremely high river flow measurements.

The function `filter()` is useful in such cases. Suppose we classified a river flow exceeding $100 \text{ m}^3/\text{s}$ at the gauge of Bathford in Section 1.1.4 as extremely high. We can then extract the subset of observations exceeding $100\text{m}^3/\text{s}$ using `filter()`, and we use `slice_head()` to print the first five observations:

```
Bathford_RF_High <- filter( Bathford_RF, RiverFlow > 100 )
slice_head( Bathford_RF_High, n=5 )
```

	Date	RiverFlow
1	1970-11-19	104.3
2	1971-01-21	128.4
3	1971-01-22	114.8
4	1971-01-23	115.2
5	1971-01-24	133.1

The function `filter()` can also handle multiple conditions. For instance, we can extract the days across the period 1991-2023 when the river flow exceeded 100m³/s using

```
Bathford_RF_High <- filter( Bathford_RF, RiverFlow > 100, year(Date) > 1990 )
slice_head( Bathford_RF_High, n=5 )
```

	Date	RiverFlow
1	1991-01-10	112.1
2	1992-11-26	106.9
3	1992-11-27	100.9
4	1992-11-28	100.6
5	1992-11-29	131.6

2.2.2 Selecting variables

Not all variables in a data set may be of interest to us. For instance, meteorological data sets often provide measurements for multiple weather variables, but we may only need to analyze precipitation and temperature.

Example: Let's consider the data set "Tuscany.csv" which provides information on the population in Tuscany, Italy, for 2020:

```
Tuscany_raw <- read.csv( "Data/Tuscany.csv" )
slice_head( Tuscany_raw, n=5 )
```

	Year	Postal_Code	Town	Province	Age	Men	Women
1	2020	45001	Aulla	MS	0	32	30
2	2020	45001	Aulla	MS	1	30	34
3	2020	45001	Aulla	MS	2	43	39
4	2020	45001	Aulla	MS	3	50	35
5	2020	45001	Aulla	MS	4	38	29

Suppose we only want to compare the population data for the different provinces and towns. As such, we don't need the variables `Year`, since all data are from 2020, and `Postal_Code`. Let's look at two possible options to achieve this using the `select()` function in dplyr.

The first option is to specify the variables we want to keep

```
Tuscany <- select( Tuscany_raw, Town:Women )
slice_head( Tuscany, n=5 )
```

	Town	Province	Age	Men	Women
1	Aulla	MS	0	32	30
2	Aulla	MS	1	30	34
3	Aulla	MS	2	43	39
4	Aulla	MS	3	50	35
5	Aulla	MS	4	38	29

Here, the colon sign indicates that we want to keep all columns from `Town` to `Women`.

The second option is to specify the variables to be excluded using the minus sign,

```
Tuscany <- select( Tuscany_raw, -Year, -Postal_Code )
```

Whether we specify the variables to be kept, or the variables to be removed, really depends on the number of variables to be included (or excluded) - we want to write as little code as possible.

2.2.3 Creating and attaching new variables

When analyzing real-world data, it may be useful to create new variables which we believe to be interesting to explore. For instance, for the population from Tuscany, we may want to calculate the total population for each age group and town, and attach this information as a new variable to the data frame.

The **mutate()** function in the dplyr package is really useful in such situations, as we can produce and directly attach a new variable **Population** using

```
Tuscany <- mutate( Tuscany, Population = Men + Women )
slice_head( Tuscany, n=5 )
```

	Town	Province	Age	Men	Women	Population
1	Aulla	MS	0	32	30	62
2	Aulla	MS	1	30	34	64
3	Aulla	MS	2	43	39	82
4	Aulla	MS	3	50	35	85
5	Aulla	MS	4	38	29	67

We see that **mutate()** requires us to provide a new variable name and to define how the values of this new variable are to be derived. Note, the function **mutate()** can also be used to attach values stored in another R object to the data frame.

Important: If you use a variable name that already exists within the data frame, **mutate()** will overwrite this column with the new values - so we can also use **mutate()** to modify the columns in your data frame.

2.2.4 Combining multiple operations - the pipe

We have already introduced quite a few useful functions for data cleaning and wrangling. Let's now consider the case that we want to combine these functions. For instance, we may want to derive the population per age group and town, and then remove the variables **Year** and **Postal_Code** from the original data frame.

How can we do this?

The first option is to manipulate the data step by step and to always store the R object after finishing one operation (similar to what we have done so far). This would be implemented as

```
Tuscany <- mutate( Tuscany_raw, Population = Men + Women )
Tuscany <- select( Tuscany, -Year, -Postal_Code )
```

This is quite a bit of code, because we have to type **Tuscany** in each line.

Can we do better?

Well, we could place all the operations into a single line

```
Tuscany <- select( mutate( Tuscany_raw, Population = Men + Women ), Town:Population )
```

However, such an approach may quickly lead to a large number of brackets, which increases the risk of frustrating syntax errors - remember this may only be the start of our analysis.

Luckily, we can avoid both these two options by using the **pipe command** `%>%` in the dplyr R package. The same commands as above would be implemented as

```
Tuscany <- Tuscany_raw %>%
  mutate( Population = Men + Women ) %>%
  select( -Year, -Postal_Code )
slice_head( Tuscany, n=5 )
```

	Town	Province	Age	Men	Women	Population
1	Aulla	MS	0	32	30	62
2	Aulla	MS	1	30	34	64
3	Aulla	MS	2	43	39	82
4	Aulla	MS	3	50	35	85
5	Aulla	MS	4	38	29	67

The operations are executed from top to bottom: We take the data frame **Tuscany_raw**, then apply the `mutate()` function to create the column **Population**, and conclude by removing the columns **Year** and **Postal_Code** from the created data frame using the `select()` function.

Tip: Combining multiple R commands can be tricky at first. If you are unsure, try to outline the way you want to manipulate the data before starting to implement it in R.

2.2.5 Summarizing the data

For large data sets, we usually want to provide data summaries. For instance, one important summary for the Tuscany data set may be the total number of people within the data. In such situations, we can apply functions such as `sum()` directly

```
sum( Tuscany$Population )
```

```
[1] 3691409
```

If we want to extract several such summaries, we can either derive each summary individually, or use the `summarize()` function in the dplyr R package. Let's also extract the proportion of men and women

```
Tuscany %>%
  summarize( "Population_Tuscany_2020" = sum( Population ),
             "Men_Tuscany_2020" = sum( Men ),
             "Women_Tuscany_2020" = sum( Women ) )
```

	Population_Tuscany_2020	Men_Tuscany_2020	Women_Tuscany_2020
1	3691409	1787649	1903760

The `summarize()` function really starts to shine when we combine it with the `group_by()` function.

Suppose we wanted the population numbers for each of the provinces, which requires us to sum up the numbers across towns and age groups while accounting for the variable **Province**. We can do this using `group_by()` and `summarize()`:

```
Tuscany_Province <- Tuscany %>%
  group_by( Province ) %>%
  summarize( Total = sum(Population) )
Tuscany_Province
```

```
# A tibble: 10 x 2
  Province   Total
  <chr>     <int>
1 AR         336450
2 FI         997940
3 GR         217803
4 LI         328855
5 LU         383688
6 MS         189786
7 PI         417799
8 PO         265153
9 PT         290177
10 SI        263758
```

The `group_by()` function splits the data subject to the specified variable (**Province** in this case) and, for each subset, the `summarize()` function then derives the population total.

Remark: We can specify multiple variables in `group_by()` to define the subgroups based on several criteria.

Let's consider a slightly more complicated task. Suppose we were asked to study the age profile of women within the population. To extract the proportion of women of a certain age, we need to group women by **Age**, but also keep track of the total number of women within the population. One possible way to extract the proportions is as follows:

```
Tuscany_Women_Age <- Tuscany %>%
  group_by( Age ) %>%
  summarize( Number = sum(Women) ) %>%
  mutate( Proportion = Number / sum(Number) )
```

Note, we used the fact that the `summarize()` function returns a data frame, and thus we can perform further operations. Finally, let's illustrate the calculated proportions using a bar plot:

```
barplot( Proportion~Age, data=Tuscany_Women_Age )
```




Figure 2.2: Age profile of living women in Tuscany for the year 2020.

We see that the highest proportions are observed for ages 40-70. The lower proportions for younger ages reflect the decrease in birth rates recorded for many countries over the past years. The decreasing proportion beyond 70 is presumably due to an increased rate of mortality for these age groups.

2.2.6 Sorting the data frame based on a variable

You may have already seen the `sort()` command, which allows you to order the values within a vector. When we consider a data frame, we may want to sort its rows subject to the values in one of the columns. For instance, we may want to sort provinces based on their population.

The function `arrange()` in the dplyr R package does exactly this job,

```
arrange( Tuscany_Province, Total )
```

```
# A tibble: 10 x 2
  Province   Total
  <chr>     <int>
1 MS        189786
2 GR        217803
3 SI        263758
4 PO        265153
5 PT        290177
6 LI        328855
7 AR        336450
8 LU        383688
9 PI        417799
10 FI       997940
```

We see that “FI” (Firenze) has the highest population among the provinces in Tuscany. Further, the output demonstrates that the default setting for `arrange()` is to sort the values in ascending order. Should we want to sort values in descending order, we have to use the additional command `desc()`:

```
Tuscany_Province %>% arrange( desc(Total) )
```

Remark: If two observations have the same value, they are listed in their original order, regardless of whether we sort in ascending or descending order. If we want to change this (which we sometimes want), we can specify a second variable in `arrange()`, just as for `group_by()`.

2.3 Working with multiple data sets

So far we have focused on analyzing a single data file. In many applications, however, data is stored across multiple data files. For instance, we may have one data file containing weather data and another data file providing insurance data related to weather-related damages. In these cases, we want to combine the different data files into a single data frame for our analysis.

The `dplyr` R package provides the functions `inner_join()`, `left_join()`, `right_join()` and `full_join()` to combine data frames based on a “key”. All these functions combine two data frames and their application is illustrated via an example in Section 1.3.1.

When working with multiple data sets, we may also want to automate the process. Imagine you had weather measurements for over 100 sites - you do not really want to spend hours just to merge the data frames. This aspect is considered in Section 1.3.2.

2.3.1 Merging two data sets

In Section 1.1.4, we focused on the river flow data collected at Bathford. The National River Flow Archive provides data for another gauge located to the west of Bath city centre; you can find the data file “Bath River Flow.csv” on Moodle. Our aim is to combine the river flow measurements into a single data frame.

We start by again loading the data for Bathford and renaming the variables,

```
Bathford_RF <- read.csv( "Data/Bathford River Flow.csv", skip=20, header=FALSE,
                           colClasses = c("character","numeric","NULL") )
Bathford_RF <- rename( Bathford_RF, Date = V1, RiverFlow = V2 )
```

A closer look at data file for the Bath gauge suggests that the data format is similar to that for Bathford. The only difference is that we now have to ignore the first 19 instead of the first 20 lines:

```
Bath_RF <- read.csv( "Data/Bath River Flow.csv", skip=19, header=FALSE,
                           colClasses = c("character","numeric","NULL") )
Bath_RF <- rename( Bath_RF, Date = V1, RiverFlow = V2 )
```

Let’s investigate the first element in each data frame:

```
Bath_RF %>% slice_head( n=1 )
```

	Date	RiverFlow
1	1976-09-01	3.39

```
Bathford_RF %>% slice_head( n=1 )
```

	Date	RiverFlow
1	1969-10-27	3.998

We see that the two gauges started operating in different years - Bath in 1976 and Bathford in 1969. So the number of rows in the two data frames is different.

When combining the two data frames, we want to match observations based on the variable **Date**, this is our “key”. Here we use the function `full_join()`, which ensures that all observations for Bath and Bathford are contained in the combined data set, and we specify that observations should be matched based on the variable **Date**,

```
RF <- Bathford_RF %>% full_join( Bath_RF, by=c("Date" = "Date") )
glimpse( RF )
```

```
Rows: 19,697
Columns: 3
$ Date      <chr> "1969-10-27", "1969-10-28", "1969-10-29", "1969-10-30", "1~
$ RiverFlow.x <dbl> 3.998, 3.958, 4.210, 4.480, 4.205, 3.830, 3.723, 3.986, 4.~
$ RiverFlow.y <dbl> NA, NA~
```

We see that the values for the first dates are correctly identified as being missing for Bath - the gauge was not in operation at the time. We are left with changing the variable names and converting the data type of **Date**

```
RF$date <- as_date( RF$date, format="%Y-%m-%d" )
RF <- rename( RF, Bathford = RiverFlow.x, Bath = RiverFlow.y )
```

Let’s plot the observations for Bath and Bathford against each other,

```
plot( RF$Bath, RF$Bathford, cex.lab = 1.5, pch=19,
      xlab="River Flow at Bath", ylab="River Flow at Bathford" )
```




Figure 2.3: Comparison of river flow for Bath and Bathford for 1 September 1976 - 30 October 2023.

What can we conclude from this plot?

Remark: If we want the first element in the combined data frame to be 01/09/1976 (the date when the gauge at Bath started operations), we would use the function `inner_join()`,

```
RF_1976_2020 <- Bathford_RF %>% inner_join( Bath_RF, by=c("Date" = "Date") )
```

Remark: The function `inner_join()` does not remove the dates after 1976 for which the observations for Bath (or Bathford) are missing, but only the days which are not listed in both files.

2.3.2 Merging multiple data sets

In practice we may work with N data sets of the same (or a very similar) format. For instance, we may have 20 data sets, and each data set contains the river flow measurements for a gauge in Somerset. Then, we do not want to implement a lot of code of the form in Section 1.3.1 just to combine all these data sets into a single data frame. Instead, we will use the **for()** loop in R.

Example: Suppose that, in addition to the river flow measurements for Bath and Bathford, we also need to consider the observations for Compton Dando, a small village to the west of Bath not located at the River Avon. For our analysis, it may be good to combine all three data sets into a single data frame, and the following piece of code is one way to create it.

We start by defining the file names and the number of lines that we have to ignore when loading the data files

```
gauges <- c( "Bath", "Bathford", "Compton Dando" )
lines_to_ignore <- c( 19, 20, 20)
```

The next step is to load the data from the different files, store the data frames in a list we call **RF_individual**, and update the variable names.

```
setwd("Data/")
RF_individual <- list()
for( k in 1:length(gauges) ){

  ## Load the data from the .csv file
  file_name <- paste( gauges[k], "River Flow.csv" )
  RF_individual[[k]] <- read.csv( file_name, skip=lines_to_ignore[k], header=FALSE,
                                 colClasses = c("character","numeric","NULL") )

  ## Change the variable names
  names( RF_individual[[k]] ) <- c( "Date", gauges[k] )

}
```

The code above includes two functions you may not have used so far and so we briefly describe them:

- 1) **paste()** is used to append “River Flow.csv” to the name of the gauge to get the file name.
- 2) **names()** is used to rename the variable names. In this case, this function was easier to use than **rename()**; the latter does not like to be given names from a vector.

Now we are ready to merge the different data frames by repeatedly using the function **full_join()**:

```
RF <- RF_individual[[1]]
for( k in 2:length(gauges) )
  RF <- RF %>% full_join( RF_individual[[k]], by=c("Date"="Date") )
glimpse( RF )
```

```
Rows: 23,955
Columns: 4
$ Date          <chr> "1976-09-01", "1976-09-02", "1976-09-03", "1976-09-04"~
$ Bath          <dbl> 3.39, 2.83, 2.97, 2.81, 2.90, 2.81, 2.59, 3.11, 2.78, ~
$ Bathford      <dbl> 2.811, 2.560, 2.337, 2.385, 2.146, 2.359, 2.367, 2.416~
$ `Compton Dando` <dbl> 0.188, 0.173, 0.170, 0.172, 0.174, 0.176, 0.195~
```

The final step is to convert the type of the variable **Date** and to sort observations by date; we know that Bathford started collecting data in 1969 but the first entry is for 1976. So we obtain the final data frame using

```
RF <- RF %>%
  mutate( Date = as_date( Date, format="%Y-%m-%d" ) ) %>%
  arrange( Date )
slice_head( RF, n=5 )
```

	Date	Bath	Bathford	Compton	Dando
1	1958-03-01	NA	NA		2.97
2	1958-03-02	NA	NA		2.32
3	1958-03-03	NA	NA		1.98
4	1958-03-04	NA	NA		1.70
5	1958-03-05	NA	NA		1.42

We can now start our analysis, for instance, by plotting the different river flows against each other:

```
par( mfrow=c(1,3), cex.lab = 1.5 )
plot( RF$Bath, RF$Bathford, pch=19, xlab="Bath", ylab="Bathford" )
plot( RF$Bath, RF$`Compton Dando`, pch=19, xlab="Bath", ylab="Compton Dando" )
plot( RF$Bathford, RF$`Compton Dando`, pch=19, xlab="Bathford", ylab="Compton Dando" )
```




Figure 2.4: Scatter plots of river flow measurements for each pair of gauges in the combined data set

What do you conclude from these plots?

Remark: If we wanted to add more gauges to the data, we only need to update the first two lines of R code in this example; the rest of the code can be left unchanged.

2.4 Summary

We have covered some of the key concepts regarding data cleaning and wrangling:

- Ensure that variables have the correct type and are given informative names
- Use the dplyr R package when working with a single data frame. The package allows you to create subsets, sort the data, etc.

- In many real-world applications we have to combine multiple data sets. The dplyr R package also provides functions to achieve this.

Important: Hardly any real-world data set is “standard” - we had to use some additional functions/options for the river flow data, as well as the Airbnb data analyzed in Problem Class 1. While we introduced useful functions to perform data wrangling, we still usually have to investigate the data file “by hand” before loading the data into R. In this course we cannot possibly cover all scenarios that may occur when working with real-world data, but you can usually find a satisfying solution using Google (or other search engines).

Chapter 3

Data Visualization

In Chapter 1 we explored how to restructure a data set and extract summaries. In this chapter we will focus on **data visualization**, i.e., the creation and interpretation of plots that give us further insight into the data. Data visualization is not just important for data exploration, but also for presenting and communicating results.

There are two very important aspects we need to keep in mind:

- 1) Effective data visualization is more about clear communication than creating impressive plots. Your analysis may be excellent, but it won't attract any attention if you cannot convey your results effectively.
- 2) Plots support our arguments and/or highlight the reason for our conclusion. As such we should interpret plots in the context of the research question and not just provide a plot as the answer.

We start by introducing a general framework for describing data graphics in Section 2.1. Sections 2.2 and 2.3 then demonstrate how to use the R package **ggplot2** for data visualization. Finally, some further aspects are considered in Section 2.4.

Remark: The methods and techniques considered in this chapter cover general aspects of data visualization. Specific methods for illustrating text and spatial data are left to the next chapters.

Important: Data visualization is to some degree subjective, because there is often not just one way to visualize the data. However, you should follow the principles outlined in this chapter, and your conclusions need to be supported by your plot. Do not claim something that is not clearly visible in your output!

3.1 Background on data visualization

Before starting to create plots in R, we establish a framework to analyze plots in terms of four basic elements: **visual cues**, **coordinate system**, **scale** and **context**. Understanding these elements will help us with producing our own plots later.

3.1.1 Visual cues

Visual cues are graphical elements that draw the audience to the aspects we want them to focus on. The book “*Data points: Visualization that means something*” by Nathan Yau (link provided on Moodle) lists nine distinct visual cues to encode a category or quantity:

- Position (quantity) - relation to other things

- Length (quantity) - size in one dimension
- Angle (quantity) - width of angle may, for instance, represent proportions (pie chart)
- Direction (quantity) - slope of line
- Shape (category) - which observations are in the same group
- Area (quantity) - size in two dimensions
- Volume (quantity) - size in three dimensions
- Shade (quantity or category) - shade in comparison to others, or grouping
- Colour (quantity or category) - colour in comparison to others, or grouping

Research has shown that our ability to perceive differences in magnitude descends in this order. One of many publications supporting this argument is “*Graphical perception: Theory, experimentation, and application to the development of graphical methods.*”, which you can find on Moodle.

Important: One crucial conclusion is that we should not rely too much on colour. Many people have colour deficiencies, which makes it very hard for them to distinguish certain colours. Consequently, before using colour, we should consider whether we could use shapes or shades instead.

Remark: In this course we focus on creating 2D graphics. While 3D plots and animations allow us to visualize a larger number of variables (and you may think they look more impressive), I would avoid using such plots except for a very limited number of cases. That’s because it is often difficult to see the exact positions of the points.

3.1.2 Example: Freediving world records

The data set “Freediving Records.csv” provides information on the progression of the world record in multiple disciplines for men and women. We will focus on the discipline “dynamic apnea with fins (DYN)” and visualize how the world records for men and women have progressed over time:




Figure 3.1: Development of the world record in dynamic apnea with fins (DYN) for men and women between 1993 and 2020.

We can identify that the following visual cues have been used:

- Shape and colour indicate whether the observations refer to the world record for men or women. Note, there is no issue with using multiple cues for the same information.